



31.7.-10.8.

PETNICA SUMMER INSTITUTE

**MACHINE
LEARNING**




PyTorch Intro

Overview of PyTorch for beginner ML practitioners

Stefan Mojsilovic

Everseen



ML/DL Frameworks

- Could we write all ML code in NumPy?
- Why shouldn't we write all ML code in NumPy?

ML/DL Frameworks

- Can we write all ML/DL code in NumPy?
 - Sure!
- Why shouldn't we write all ML/DL code in NumPy?
 - Missing a lot of ML/DL specific functions
 - No GPU support
 - No Automatic differentiation

ML/DL Frameworks

- A set of tools allowing you to focus on specific tasks
- A skeleton upon which you build an application
- Takes care of the low level stuff
- Allows you to concentrate on what matters the most

ML/DL Frameworks



PyTorch vs Tensorflow




Image source: <https://thegradient.pub/>

PyTorch vs other frameworks

- **PyTorch** (Facebook, IBM, Yandex, OpenAI) dominates research applications
- **Tensorflow** (Google, Uber, ABnB) and **MXNET** (Amazon, Microsoft, Intel) dominate industry applications
 - Exception is caffe2 (now part of pytorch) which is Facebook's default production environment
- Useful links:
 - <https://www.netguru.com/blog/deep-learning-frameworks-comparison>
 - <https://thegradient.pub/state-of-ml-frameworks-2019-pytorch-dominates-research-tensorflow-dominates-industry/>
 - <https://executecommands.com/machine-learning-frameworks-2020/>

What is PyTorch?

Python-based package for scientific computing with two goals:

- A replacement for NumPy to use the power of GPUs
- A deep learning research platform w/ maximum flexibility and speed

Why is it good?

- Intuitive, low-level, highly modular, easy to extend, open source
- High growth – numerous libraries and support by the devs
- Massively adopted by the research community

PyTorch as Numpy replacement

- torch.Tensor

Numpy	PyTorch
<code>x = numpy.array(object, dtype=None, ...)</code>	<code>x = torch.tensor(data, dtype=None, device=None, requires_grad=False)</code>
<code>.ones()</code> <code>.zeros()</code>	<code>.ones()</code> <code>.zeros()</code>
<code>.max()</code> <code>.mean()</code> <code>.sum()</code>	<code>.max()</code> <code>.mean()</code> <code>.sum()</code>
<code>.random.rand()</code> <code>.random.normal()</code> <code>.random.randint()</code>	<code>.rand()</code> <code>.normal()</code> <code>.randint()</code>
<code>x.shape</code> <code>x.ndim</code> <code>len(x)</code> <code>x.size</code> <code>x.dtype</code> <code>x[0, :, 1:3]</code>	<code>x.shape</code> <code>x.ndim</code> <code>len(x)</code> <code>x.size()</code> <code>x.dtype</code> <code>x[0, :, 1:3]</code>
<code>+</code> <code>-</code> <code>*</code> <code>/</code>	<code>+</code> <code>-</code> <code>*</code> <code>/</code>
<code>.matmul(x, y)</code>	<code>.mm(x, y)</code>
<code>x.reshape()</code> <code>x.T</code> <code>x.copy()</code>	<code>x.reshape()</code> <code>x.view()</code> <code>x.t()</code> <code>x.clone()</code>
<code>x.squeeze()</code> <code>x.expand_dims()</code>	<code>x.squeeze()</code> <code>x.unsqueeze</code>
<code>np.save('file.npy', x)</code> <code>np.load('file.npy')</code>	<code>torch.save(x, 'file.pth')</code> <code>torch.load('file.pth')</code>

From CPU to GPU and back

1. `x = x.to(device="cuda")`

2. `x = x.to(device="cpu")`

Literally as simple as that 😊

- Arithmetic and torch functions on GPU tensors will be done on GPU
- Note: mixing operands with different devices raises a RuntimeError!

PyTorch as a DL Research Platform

- torch.cuda -> tools for enabling GPU computation
- torch.autograd -> tools for automatic differentiation
- torch.data -> tools for loading your data efficiently
- torch.nn -> tools for building your ML models
- torch.optim -> tools for optimizing your models
- Libraries -> torchaudio, torchtext, torchvision...
- And many, many, many more:
 - <https://pytorch.org/ecosystem/>
 - <https://pytorch.org/hub/>

torch.cuda

- Info on CUDA-enabled devices

- `.is_available()`
- `.device_count()`
- `.current_device()`
- `.get_device_name()`

- Selecting a device

- `device = torch.device("cuda:n")`, where `n=0,1,2...`
- `.set_device(device)`

- Advanced: Explicit control over streams, memory, events, sync...

```
1 print(torch.cuda.is_available())
2 print(torch.cuda.device_count())
3 print(torch.cuda.current_device())
4 print(torch.cuda.get_device_name())
```

True

1

0

GeForce GTX 1060

torch.autograd

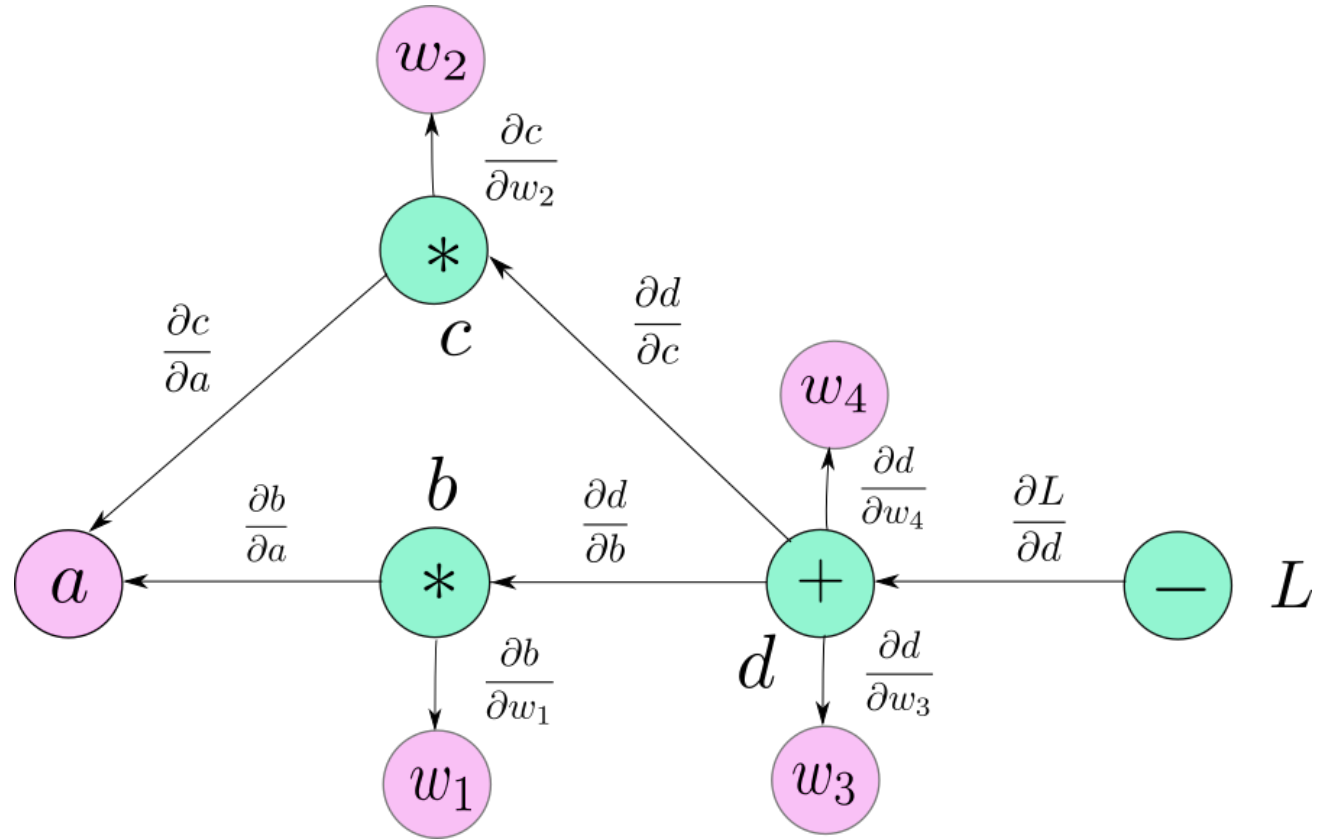
Dynamic Computational Graph

$$b = w_1 * a$$

$$c = w_2 * a$$

$$d = w_3 * b + w_4 * c$$

$$L = 10 - d$$



torch.autograd

- .Function
 - .forward()
 - .backward()
- Tensor attributes and methods
 - .requires_grad
 - .is_leaf
 - .grad_fn
 - .grad
 - .detach()
- Backprop with .backward()
 - On scalar valued tensors only
 - Frees the non-leaf buffers and destroys the graph
 - Remember to zero the gradients between .backward() calls, otherwise they will be accumulated

Note:

- To turn off requires_grad globally:
torch.set_grad_enabled(False)
- Or as context manager - with torch.no_grad():

torch.utils.data

- .Dataset -> feat, label = dataset[idx]
 - .__init__()
 - .__len__()
 - .__getitem__(idx)
- .DataLoader -> for feat, label in dl ...
 - Shuffling
 - Batching
 - Prefetching
 - Multiprocessing

```
1 class ImageDataset(Dataset):
2     """Image dataset for classification."""
3
4     def __init__(self, root_dir, csv_file, transform=None):
5         """
6         Args:
7             root_dir (string): Directory with all the images.
8             csv_file (string): Path to the csv file with paths to images
9                               relative to the root_dir, and their class labels.
10            transform (callable): Optional transform to be applied on a sample.
11        """
12        self.root_dir = root_dir
13        self.image_dataframe = pd.read_csv(csv_file)
14        self.transform = transform
15
16    def __len__(self):
17        return len(self.image_dataframe)
18
19    def __getitem__(self, idx):
20        if torch.is_tensor(idx):
21            idx = idx.tolist()
22
23        img_filepath = os.path.join(self.root_dir,
24                                    self.image_dataframe.iloc[idx, 0])
25        image = cv2.imread(img_filepath)
26        class_label = self.image_dataframe.iloc[idx, 1]
27
28        if self.transform:
29            image = self.transform(image)
30
31        return image, label
```


torch.nn.Module

- Used for building NNs and their components
 - Layers (fc, conv, pooling, normalization, dropout...)
 - Activations (sigmoid, softmax, ReLU...)
 - Losses (BCE, CE, MSE, triplet)
- Or as a container for already built components
 - Manually adding them to the Module
 - Or using nn.Sequential

```
1 model = nn.Sequential(  
2     nn.Linear(100, 10),  
3     nn.ReLU(),  
4     nn.Linear(10, 1),  
5     nn.Sigmoid() )
```

torch.nn.Module

- When defining a model, subclass Module and implement:
 - `__init__()` – define the parameters of the object
 - `forward()` – define the computational graph

```
1 import torch.nn as nn
2 import torch.nn.functional as F
3 class Model(torch.nn.Module):
4     def __init__(self):
5         super(Model, self).__init__()
6         self.layer1 = torch.nn.Linear(in_features=100, out_features=10, bias=True)
7         self.layer2 = torch.nn.Linear(in_features=10, out_features=1, bias=True)
8
9     def forward(self, x):
10         x = F.relu(self.layer1(x))
11         x = self.layer2(x)
12         return x
```

- However, when calling the instance of your model class
 - USE: `output = model(input)`
 - NOT: `output = model.forward(input)`, as it will disregard hooks

torch.optim

- Optimizers – optimization algorithm (SGD, Adam...)
 - `optimizer=optim.NAME(model.parameters(), **kwargs)`
 - `optimizer.zero_grad()`
 - ... model output, loss, `loss.backward()`
 - `optimizer.step()` – uses gradients computed by `.backward()` to update `model.parameters()`
- LR Schedulers – update optimizer hyper-parameters
 - `lr_scheduler=optim.lr_scheduler.NAME(optimizer, **kwargs)`
 - `lr_scheduler.step()` – updates the optimizer hyper-parameters

Congratulations!



PyTorch Intro

By Stefan Mojsilovic

stefan.mjs+psiml@gmail.com

Working in Everseen Ltd. on deep learning applications in computer vision and large-scale retail. Interested in cognitive models and augmented intelligence.

